

AE 04: NYC flights + data wrangling

Gracie Carlaw

```
library(tidyverse)
library(nycflights13)
```

Exercise 1

Your turn: Fill in the blanks:

The `flights` data frame has 336776 rows. Each row represents a different flight.

Exercise 2

Your turn: What are the names of the variables in `flights`.

```
names(flights)
```

[1]	"year"	"month"	"day"	"dep_time"
[5]	"sched_dep_time"	"dep_delay"	"arr_time"	"sched_arr_time"
[9]	"arr_delay"	"carrier"	"flight"	"tailnum"
[13]	"origin"	"dest"	"air_time"	"distance"
[17]	"hour"	"minute"	"time_hour"	

Exercise 3 - `select()`

- Make a data frame that only contains the variables `dep_delay` and `arr_delay`.

```
select(flights, "dep_delay", "arr_delay")
```

```
# A tibble: 336,776 x 2
```

```
  dep_delay arr_delay
    <dbl>      <dbl>
1         2         11
2         4         20
3         2         33
4        -1        -18
5        -6        -25
6        -4         12
7        -5         19
8        -3        -14
9        -3         -8
10       -2          8
```

```
# i 336,766 more rows
```

- Make a data frame that keeps every variable except `dep_delay`.

```
select(flights, -dep_delay)
```

```
# A tibble: 336,776 x 18
```

```
  year month   day dep_time sched_dep_time arr_time sched_arr_time arr_delay
  <int> <int> <int>   <int>         <int>      <int>         <int>      <dbl>
1  2013     1     1     517           515         830           819         11
2  2013     1     1     533           529         850           830         20
3  2013     1     1     542           540         923           850         33
4  2013     1     1     544           545        1004          1022        -18
5  2013     1     1     554           600         812           837        -25
6  2013     1     1     554           558         740           728         12
7  2013     1     1     555           600         913           854         19
8  2013     1     1     557           600         709           723        -14
9  2013     1     1     557           600         838           846         -8
10 2013     1     1     558           600         753           745          8
```

```
# i 336,766 more rows
```

```
# i 10 more variables: carrier <chr>, flight <int>, tailnum <chr>,
#   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>
```

- Make a data frame that includes all variables between `year` through `dep_delay` (inclusive). These are all variables that provide information about the departure of each flight.

```
select(flights, year:dep_delay)
```

```
# A tibble: 336,776 x 6
  year month   day dep_time sched_dep_time dep_delay
  <int> <int> <int>   <int>         <int>         <dbl>
1  2013     1     1     517             515           2
2  2013     1     1     533             529           4
3  2013     1     1     542             540           2
4  2013     1     1     544             545          -1
5  2013     1     1     554             600          -6
6  2013     1     1     554             558          -4
7  2013     1     1     555             600          -5
8  2013     1     1     557             600          -3
9  2013     1     1     557             600          -3
10 2013     1     1     558             600          -2
# i 336,766 more rows
```

- Use the `select` helper `contains()` to make a data frame that includes the variables associated with the arrival, i.e., contains the string `"arr_"` in the name.

```
flights %>%
  select(contains("arr"))
```

```
# A tibble: 336,776 x 4
  arr_time sched_arr_time arr_delay carrier
  <int>         <int>         <dbl> <chr>
1     830           819           11 UA
2     850           830           20 UA
3     923           850           33 AA
4    1004          1022          -18 B6
5     812           837          -25 DL
6     740           728           12 UA
7     913           854           19 B6
8     709           723          -14 EV
9     838           846           -8 B6
10    753           745            8 AA
# i 336,766 more rows
```

Exercise 4 - `slice()`

- Display the first five rows of the `flights` data frame.

```
flights %>%
  slice_head(n = 5)
```

```
# A tibble: 5 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517           515           2     830           819
2  2013     1     1     533           529           4     850           830
3  2013     1     1     542           540           2     923           850
4  2013     1     1     544           545          -1    1004          1022
5  2013     1     1     554           600          -6     812           837
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Display the last two rows of the flights data frame.

```
flights %>%
  slice_tail(n = 2)
```

```
# A tibble: 2 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     9    30      NA           1159          NA      NA           1344
2  2013     9    30      NA           840          NA      NA           1020
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Exercise 5 - arrange()

- Let's arrange the data by departure delay, so the flights with the shortest departure delays will be at the top of the data frame.

```
flights %>%
  arrange(dep_delay)
```

```
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     12     7     2040             2123         -43     40           2352
2  2013      2     3     2022             2055         -33    2240           2338
3  2013     11    10     1408             1440         -32    1549           1559
4  2013      1    11     1900             1930         -30    2233           2243
5  2013      1    29     1703             1730         -27    1947           1957
6  2013      8     9      729             755         -26    1002            955
7  2013     10    23     1907             1932         -25    2143           2143
8  2013      3    30     2030             2055         -25    2213           2250
9  2013      3     2     1431             1455         -24    1601           1631
10 2013      5     5      934             958         -24    1225           1309
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Question: What does it mean for the `dep_delay` to have a negative value?

It likely means that the flight departed early.

- Arrange the data by descending departure delay, so the flights with the longest departure delays will be at the top.

```
flights %>%
  arrange(desc(dep_delay))
```

```
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013      1     9      641             900        1301    1242           1530
2  2013      6    15     1432             1935        1137    1607           2120
3  2013      1    10     1121             1635        1126    1239           1810
4  2013      9    20     1139             1845        1014    1457           2210
5  2013      7    22      845             1600        1005    1044           1815
6  2013      4    10     1100             1900         960    1342           2211
7  2013      3    17     2321             810         911     135           1020
8  2013      6    27      959             1900         899    1236           2226
9  2013      7    22     2257             759         898     121           1026
10 2013     12     5      756             1700         896    1058           2020
# i 336,766 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
```

```
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

- **Your turn:** Create a data frame that only includes the plane tail number (`tailnum`), carrier (`carrier`), and departure delay for the flight with the longest departure delay. What is the plane tail number (`tailnum`) for this flight?

```
flights %>%
  select("tailnum", "carrier") %>%
  arrange(desc("dep_delay")) %>%
  slice_head(n = 1)
```

```
# A tibble: 1 x 2
  tailnum carrier
  <chr>    <chr>
1 N14228  UA
```

The tail number for this flight is N14228.

Exercise 6 - filter()

- Filter for all rows where the destination airport is RDU.

```
flights%>%
  filter(dest == "RDU")
```

```
# A tibble: 8,163 x 19
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     800             810          -10     949             955
2  2013     1     1     832             840           -8    1006            1030
3  2013     1     1     851             851           0    1032            1036
4  2013     1     1     917             920           -3    1052            1108
5  2013     1     1    1024            1030           -6    1204            1215
6  2013     1     1    1127            1129           -2    1303            1309
7  2013     1     1    1157            1205           -8    1342            1345
8  2013     1     1    1240            1235            5    1415            1415
9  2013     1     1    1317            1325           -8    1454            1505
10 2013     1     1    1449            1450           -1    1651            1640
# i 8,153 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
```

```
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dtm>
```

- Filter for all rows where the destination airport is RDU and the arrival delay is less than 0.

```
flights %>%
  filter(dest == "RDU" & arr_delay < 0)
```

```
# A tibble: 4,232 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	800	810	-10	949	955
2	2013	1	1	832	840	-8	1006	1030
3	2013	1	1	851	851	0	1032	1036
4	2013	1	1	917	920	-3	1052	1108
5	2013	1	1	1024	1030	-6	1204	1215
6	2013	1	1	1127	1129	-2	1303	1309
7	2013	1	1	1157	1205	-8	1342	1345
8	2013	1	1	1317	1325	-8	1454	1505
9	2013	1	1	1505	1510	-5	1654	1655
10	2013	1	1	1800	1800	0	1945	1951

```
# i 4,222 more rows
```

```
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dtm>
```

- **Your turn:** Describe what the code is doing in words.

The code below is finding rows (flights) within the “flights” dataset where the destination is either “RDU” or “GSO” and the arrival and departure delays are both less than 0.

```
flights |>
  filter(
    dest %in% c("RDU", "GSO"),
    arr_delay < 0 | dep_delay < 0
  )
```

```
# A tibble: 6,203 x 19
```

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
1	2013	1	1	800	810	-10	949	955

```

2 2013      1      1      832          840          -8      1006          1030
3 2013      1      1      851          851           0      1032          1036
4 2013      1      1      917          920          -3      1052          1108
5 2013      1      1     1024         1030          -6      1204          1215
6 2013      1      1     1127         1129          -2      1303          1309
7 2013      1      1     1157         1205          -8      1342          1345
8 2013      1      1     1317         1325          -8      1454          1505
9 2013      1      1     1449         1450          -1      1651          1640
10 2013     1      1     1505         1510          -5      1654          1655
# i 6,193 more rows
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>

```

Hint: Logical operators in R:

operator	definition
<	is less than?
<=	is less than or equal to?
>	is greater than?
>=	is greater than or equal to?
==	is exactly equal to?
!=	is not equal to?
x & y	is x AND y?
x \ y	is x OR y?
is.na(x)	is x NA?
!is.na(x)	is x not NA?
x %in% y	is x in y?
!(x %in% y)	is x not in y?
!x	is not x? (only makes sense if x is TRUE or FALSE)

Exercise 7 - count()

- Create a frequency table of the destination locations for flights from New York.

```

flights %>%
  count(origin, dest, sort = TRUE)

```

```

# A tibble: 224 x 3
  origin dest      n

```



```

      <chr>  <chr> <int>
1  JFK     LAX   11262
2  LGA     ATL   10263
3  LGA     ORD    8857
4  JFK     SFO    8204
5  LGA     CLT    6168
6  EWR     ORD    6100
7  JFK     BOS    5898
8  LGA     MIA    5781
9  JFK     MCO    5464
10 EWR     BOS    5327
# i 214 more rows

```

- In which month was there the fewest number of flights? How many flights were there in that month?

```

flights %>%
  count(month, sort = TRUE)

```

```

# A tibble: 12 x 2
  month     n
  <int> <int>
1     7 29425
2     8 29327
3    10 28889
4     3 28834
5     5 28796
6     4 28330
7     6 28243
8    12 28135
9     9 27574
10    11 27268
11     1 27004
12     2 24951

```

The fewest flights were in February with 24951 flights.

- **Your turn:** On which date (month + day) was there the largest number of flights? How many flights were there on that day?

```

flights %>%
  count(month, day, sort = TRUE)

```

```
# A tibble: 365 x 3
  month   day     n
  <int> <int> <int>
1     11    27 1014
2      7    11 1006
3      7     8 1004
4      7    10 1004
5     12     2 1004
6      7    18 1003
7      7    25 1003
8      7    12 1002
9      7     9 1001
10     7    17 1001
# i 355 more rows
```

The day with the most flights was November 27th with 1014 flights.

Exercise 8 - mutate()

- Convert `air_time` (minutes in the air) to hours and then create a new variable, `mph`, the miles per hour of the flight.

```
flights %>%
  mutate(air_time / 60,
         mph = distance/air_time)
```

```
# A tibble: 336,776 x 21
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
1  2013     1     1     517           515           2     830           819
2  2013     1     1     533           529           4     850           830
3  2013     1     1     542           540           2     923           850
4  2013     1     1     544           545          -1    1004          1022
5  2013     1     1     554           600          -6     812           837
6  2013     1     1     554           558          -4     740           728
7  2013     1     1     555           600          -5     913           854
8  2013     1     1     557           600          -3     709           723
9  2013     1     1     557           600          -3     838           846
10 2013     1     1     558           600          -2     753           745
# i 336,766 more rows
# i 13 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
```

```
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dtm>, `air_time/60` <dbl>, mph <dbl>
```

- **Your turn:** First, count the number of flights each month, and then calculate the proportion of flights in each month. What proportion of flights take place in July?

```
flights %>%
  count(month, sort = TRUE) %>%
  mutate(month_prop = n/sum(n)) %>%
  filter(month == 7)
```

```
# A tibble: 1 x 3
  month      n month_prop
  <int> <int>      <dbl>
1     7 29425      0.0874
```

The proportion of flights that take place in July is 0.08737262, or 8.73%

- Create a new variable, `rdu_bound`, which indicates whether the flight is to RDU or not. Then, for each departure airport (`origin`), calculate what proportion of flights originating from that airport are to RDU.

```
flights %>%
  mutate(rdu_bound = if_else(dest == "RDU", "YES", "NO")) %>%
  count(rdu_bound, origin, sort = TRUE) %>%
  filter(rdu_bound == "YES") %>%
  mutate(prop_rdu_bound = n/sum(n))
```

```
# A tibble: 3 x 4
  rdu_bound origin      n prop_rdu_bound
  <chr>      <chr> <int>      <dbl>
1 YES      LGA     3581      0.439
2 YES      JFK     3100      0.380
3 YES      EWR     1482      0.182
```

Exercise 9 - summarize()

- Find mean arrival delay for all flights.

```
flights %>%
  summarize(
    avg_delay = mean(arr_delay, na.rm = TRUE))
```

```
# A tibble: 1 x 1
  avg_delay
  <dbl>
1      6.90
```

Exercise 10 - group_by()

- Find mean arrival delay for for each month.

```
flights %>%
  group_by(month) %>%
  summarize(
    avg_delay = mean(arr_delay, na.rm = TRUE))
```

```
# A tibble: 12 x 2
  month avg_delay
  <int>   <dbl>
1     1     6.13
2     2     5.61
3     3     5.81
4     4    11.2
5     5     3.52
6     6    16.5
7     7    16.7
8     8     6.04
9     9    -4.02
10    10    -0.167
11    11     0.461
12    12    14.9
```

- **Your turn:** What is the median departure delay for each airports around NYC (origin)? Which airport has the shortest median departure delay?

```
flights %>%
  group_by(origin) %>%
  summarize(
    avg_dep_delay = median(dep_delay, na.rm = TRUE)) %>%
  slice_min(avg_dep_delay, n = 1)
```

```
# A tibble: 1 x 2
  origin avg_dep_delay
  <chr>      <dbl>
1 LGA         -3
```

Additional Practice

Try these on your own, either in class if you finish early, or after class.

1. Create a new dataset that only contains flights that do not have a missing departure time. Include the columns `year`, `month`, `day`, `dep_time`, `dep_delay`, and `dep_delay_hours` (the departure delay in hours). *Hint: Note you may need to use `mutate()` to make one or more of these variables.*

```
flights %>%
  mutate(dep_delay_hours = dep_delay/60) %>%
  select(year, month, day, dep_time, dep_delay, dep_delay_hours) %>%
  filter(!is.na(dep_time))
```

```
# A tibble: 328,521 x 6
  year month   day dep_time dep_delay dep_delay_hours
  <int> <int> <int>   <int>      <dbl>         <dbl>
1  2013     1     1     517         2         0.0333
2  2013     1     1     533         4         0.0667
3  2013     1     1     542         2         0.0333
4  2013     1     1     544        -1        -0.0167
5  2013     1     1     554        -6        -0.1
6  2013     1     1     554        -4        -0.0667
7  2013     1     1     555        -5        -0.0833
8  2013     1     1     557        -3        -0.05
9  2013     1     1     557        -3        -0.05
10 2013     1     1     558        -2        -0.0333
# i 328,511 more rows
```

2. For each airplane (uniquely identified by `tailnum`), use a `group_by()` paired with `summarize()` to find the sample size, mean, and standard deviation of flight distances. Then include only the top 5 and bottom 5 airplanes in terms of mean distance traveled per flight in the final data frame.

```
flights %>%
  group_by(tailnum) %>%
  summarize(mean = mean(distance),
            standard_deviation = sd(distance, na.rm = TRUE),
            n=n()) %>%
  arrange(mean) %>%
  filter(row_number() > max(row_number()) - 5 | row_number() <= 5)
```

```
# A tibble: 10 x 4
  tailnum mean standard_deviation    n
  <chr>   <dbl>          <dbl> <int>
1 N955UW  173.             32.9   225
2 N948UW  174.             32.7   232
3 N959UW  174.             34.3   213
4 N956UW  174.             31.4   222
5 N945UW  176.             31.2   285
6 N390HA  4983              0      20
7 N391HA  4983              0      21
8 N392HA  4983              0      13
9 N393HA  4983              0      10
10 N395HA  4983              0       7
```

NOTE: The last filter to have top 5 and bottom 5 row numbers I found from an online forum (<https://stackoverflow.com/questions/56809476/dplyr-filter-top-and-bottom-rows-by-value-simultaneously-on-grouped-data>). I struggled to find an easier way to do it, but that was the only thing I could find that worked. I just wanted to acknowledge that that line was not my doing.