

# AE 04: NYC flights + data wrangling

Gracie Carlaw

```
library(tidyverse)
library(nycflights13)
```

## Exercise 1

**Your turn:** Fill in the blanks:

The `flights` data frame has 336776 rows. Each row represents a different flight.

## Exercise 2

**Your turn:** What are the names of the variables in `flights`.

```
names(flights)
```

[1]	"year"	"month"	"day"	"dep_time"
[5]	"sched_dep_time"	"dep_delay"	"arr_time"	"sched_arr_time"
[9]	"arr_delay"	"carrier"	"flight"	"tailnum"
[13]	"origin"	"dest"	"air_time"	"distance"
[17]	"hour"	"minute"	"time_hour"	

## Exercise 3 - `select()`

- Make a data frame that only contains the variables `dep_delay` and `arr_delay`.

```
select(flights, "dep_delay", "arr_delay")
```

```
# A tibble: 336,776 x 2
```

```
  dep_delay arr_delay
    <dbl>      <dbl>
1         2         11
2         4         20
3         2         33
4        -1        -18
5        -6        -25
6        -4         12
7        -5         19
8        -3        -14
9        -3         -8
10       -2          8
```

```
# i 336,766 more rows
```

- Make a data frame that keeps every variable except `dep_delay`.

```
select(flights, -dep_delay)
```

```
# A tibble: 336,776 x 18
```

```
  year month   day dep_time sched_dep_time arr_time sched_arr_time arr_delay
  <int> <int> <int>   <int>         <int>      <int>         <int>      <dbl>
1  2013     1     1     517           515         830           819         11
2  2013     1     1     533           529         850           830         20
3  2013     1     1     542           540         923           850         33
4  2013     1     1     544           545        1004          1022        -18
5  2013     1     1     554           600         812           837        -25
6  2013     1     1     554           558         740           728         12
7  2013     1     1     555           600         913           854         19
8  2013     1     1     557           600         709           723        -14
9  2013     1     1     557           600         838           846         -8
10 2013     1     1     558           600         753           745          8
```

```
# i 336,766 more rows
```

```
# i 10 more variables: carrier <chr>, flight <int>, tailnum <chr>,
#   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
#   minute <dbl>, time_hour <dtm>
```

- Make a data frame that includes all variables between `year` through `dep_delay` (inclusive). These are all variables that provide information about the departure of each flight.

```
select(flights, year:dep_delay)
```

```
# A tibble: 336,776 x 6
  year month   day dep_time sched_dep_time dep_delay
  <int> <int> <int>   <int>         <int>      <dbl>
1  2013     1     1     517             515         2
2  2013     1     1     533             529         4
3  2013     1     1     542             540         2
4  2013     1     1     544             545        -1
5  2013     1     1     554             600        -6
6  2013     1     1     554             558        -4
7  2013     1     1     555             600        -5
8  2013     1     1     557             600        -3
9  2013     1     1     557             600        -3
10 2013     1     1     558             600        -2
# i 336,766 more rows
```

- Use the `select` helper `contains()` to make a data frame that includes the variables associated with the arrival, i.e., contains the string `"arr\_"` in the name.

```
flights %>%
  select(contains("arr"))
```

```
# A tibble: 336,776 x 4
  arr_time sched_arr_time arr_delay carrier
  <int>         <int>      <dbl> <chr>
1     830           819         11 UA
2     850           830         20 UA
3     923           850         33 AA
4    1004          1022        -18 B6
5     812           837        -25 DL
6     740           728         12 UA
7     913           854         19 B6
8     709           723        -14 EV
9     838           846         -8 B6
10    753           745          8 AA
# i 336,766 more rows
```

#### Exercise 4 - `slice()`

- Display the first five rows of the `flights` data frame.

```
# add code here
```

- Display the last two rows of the `flights` data frame.

```
# add code here
```

### Exercise 5 - `arrange()`

- Let's arrange the data by departure delay, so the flights with the shortest departure delays will be at the top of the data frame.

```
# add code here
```

- Question: What does it mean for the `dep_delay` to have a negative value?

Add your response here.

- Arrange the data by descending departure delay, so the flights with the longest departure delays will be at the top.

```
# add code here
```

- **Your turn:** Create a data frame that only includes the plane tail number (`tailnum`), carrier (`carrier`), and departure delay for the flight with the longest departure delay. What is the plane tail number (`tailnum`) for this flight?

```
# add code here
```

### Exercise 6 - `filter()`

- Filter for all rows where the destination airport is RDU.

```
# add code here
```

- Filter for all rows where the destination airport is RDU and the arrival delay is less than 0.

```
# add code here
```

- **Your turn:** Describe what the code is doing in words.

Add response here.

```
flights |>
  filter(
    dest %in% c("RDU", "GSO"),
    arr_delay < 0 | dep_delay < 0
  )
```

```
# A tibble: 6,203 x 19
```

```
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
1  2013     1     1     800           810        -10     949           955
2  2013     1     1     832           840         -8    1006          1030
3  2013     1     1     851           851          0    1032          1036
4  2013     1     1     917           920         -3    1052          1108
5  2013     1     1    1024          1030         -6    1204          1215
6  2013     1     1    1127          1129         -2    1303          1309
7  2013     1     1    1157          1205         -8    1342          1345
8  2013     1     1    1317          1325         -8    1454          1505
9  2013     1     1    1449          1450         -1    1651          1640
10 2013     1     1    1505          1510         -5    1654          1655
```

```
# i 6,193 more rows
```

```
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
```

**Hint:** Logical operators in R:

operator	definition
<	is less than?
<=	is less than or equal to?
>	is greater than?
>=	is greater than or equal to?
==	is exactly equal to?
!=	is not equal to?
x & y	is x AND y?
x \   y	is x OR y?
is.na(x)	is x NA?
!is.na(x)	is x not NA?
x %in% y	is x in y?
!(x %in% y)	is x not in y?
!x	is not x? (only makes sense if x is TRUE or FALSE)

### Exercise 7 - count()

- Create a frequency table of the destination locations for flights from New York.

```
# add code here
```

- In which month was there the fewest number of flights? How many flights were there in that month?

```
# add code here
```

- **Your turn:** On which date (month + day) was there the largest number of flights? How many flights were there on that day?

```
# add code here
```

### Exercise 8 - mutate()

- Convert `air_time` (minutes in the air) to hours and then create a new variable, `mph`, the miles per hour of the flight.

```
# add code here
```

- **Your turn:** First, count the number of flights each month, and then calculate the proportion of flights in each month. What proportion of flights take place in July?

```
# add code here
```

- Create a new variable, `rdu_bound`, which indicates whether the flight is to RDU or not. Then, for each departure airport (`origin`), calculate what proportion of flights originating from that airport are to RDU.

```
# add code here
```

### Exercise 9 - summarize()

- Find mean arrival delay for all flights.

```
# add code here
```

## Exercise 10 - `group_by()`

- Find mean arrival delay for for each month.

```
# add code here
```

- **Your turn:** What is the median departure delay for each airports around NYC (`origin`)? Which airport has the shortest median departure delay?

```
# add code here
```

## Additional Practice

Try these on your own, either in class if you finish early, or after class.

1. Create a new dataset that only contains flights that do not have a missing departure time. Include the columns `year`, `month`, `day`, `dep_time`, `dep_delay`, and `dep_delay_hours` (the departure delay in hours). *Hint: Note you may need to use `mutate()` to make one or more of these variables.*

```
# add code here
```

2. For each airplane (uniquely identified by `tailnum`), use a `group_by()` paired with `summarize()` to find the sample size, mean, and standard deviation of flight distances. Then include only the top 5 and bottom 5 airplanes in terms of mean distance traveled per flight in the final data frame.

```
# add code here
```